

o9s4gt6j1

November 27, 2025

1 EC5303 Team Project: BTC & ETH Dependence and Regression

Full pipeline: load Kaggle data, clean, compute returns, run distribution tests, fit copulas, and estimate regressions as specified in README.

```
[ ]: # Run once in a fresh environment to install all dependencies
# !pip install -r requirements.txt
```

```
[1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

from pathlib import Path
from scipy import stats
from statsmodels.distributions.empirical_distribution import ECDF
import statsmodels.api as sm
import kagglehub
import pyvocopilib as pv

sns.set_theme(style="whitegrid")
pd.set_option("display.float_format", lambda x: f"{x:.6f}")
```

```
d:\miniconda\envs\ec5303\Lib\site-packages\tqdm\auto.py:21: TqdmWarning:
IPProgress not found. Please update jupyter and ipywidgets. See
https://ipywidgets.readthedocs.io/en/stable/user_install.html
    from .autonotebook import tqdm as notebook_tqdm
```

1.1 Section A — Data Loading and Cleaning

- Uses `kagglehub.dataset_load` to retrieve BTC and ETH price history (Kaggle dataset `sudalairajkumar/cryptocurrencypricehistory`).
- Cleans columns, converts dates, sorts ascending, and drops duplicates.
- Computes daily log returns and merges into `df_ret` with `btc_ret` and `eth_ret`.

```
[5]: DATASET_ID = "sudalairajkumar/cryptocurrencypricehistory"
BTC_FILE = "coin_Bitcoin.csv"
```

```

ETH_FILE = "coin_Ethereum.csv"

def load_kaggle_csv(dataset_id: str, filename: str) -> pd.DataFrame:
    """Download dataset locally and read a specific CSV file."""
    dataset_dir = kagglehub.dataset_download(dataset_id)
    csv_path = Path(dataset_dir) / filename
    return pd.read_csv(csv_path)

df_btc_raw = load_kaggle_csv(DATASET_ID, BTC_FILE)
df_eth_raw = load_kaggle_csv(DATASET_ID, ETH_FILE)

display(df_btc_raw.head())
display(df_eth_raw.head())

```

Downloading from https://www.kaggle.com/api/v1/datasets/download/sudalairajkumar/cryptocurrencypricehistory?dataset_version_number=3...

100% | 1.70M/1.70M [00:01<00:00, 1.56MB/s]

Extracting files...

SNo	Name	Symbol	Date	High	Low	Open	\
0	1	Bitcoin	BTC	2013-04-29 23:59:59	147.488007	134.000000	134.444000
1	2	Bitcoin	BTC	2013-04-30 23:59:59	146.929993	134.050003	144.000000
2	3	Bitcoin	BTC	2013-05-01 23:59:59	139.889999	107.720001	139.000000
3	4	Bitcoin	BTC	2013-05-02 23:59:59	125.599998	92.281898	116.379997
4	5	Bitcoin	BTC	2013-05-03 23:59:59	108.127998	79.099998	106.250000

	Close	Volume	Marketcap
0	144.539993	0.000000	1,603,768,864.500000
1	139.000000	0.000000	1,542,813,125.000000
2	116.989998	0.000000	1,298,954,593.750000
3	105.209999	0.000000	1,168,517,495.250000
4	97.750000	0.000000	1,085,995,168.750000

SNo	Name	Symbol	Date	High	Low	Open	\
0	1	Ethereum	ETH	2015-08-08 23:59:59	2.798810	0.714725	2.793760
1	2	Ethereum	ETH	2015-08-09 23:59:59	0.879810	0.629191	0.706136
2	3	Ethereum	ETH	2015-08-10 23:59:59	0.729854	0.636546	0.713989
3	4	Ethereum	ETH	2015-08-11 23:59:59	1.131410	0.663235	0.708087
4	5	Ethereum	ETH	2015-08-12 23:59:59	1.289940	0.883608	1.058750

	Close	Volume	Marketcap
0	0.753325	674,188.000000	45,486,894.240800
1	0.701897	532,170.000000	42,399,573.499100
2	0.708448	405,283.000000	42,818,364.394500
3	1.067860	1,463,100.000000	64,569,288.432800
4	1.217440	2,150,620.000000	73,645,010.986300

```
[6]: KEEP_COLS = ["Date", "Open", "High", "Low", "Close", "Volume", "Marketcap"]

def clean_crypto_df(df: pd.DataFrame) -> pd.DataFrame:
    df = df.copy()
    df["Date"] = pd.to_datetime(df["Date"])
    df = df.sort_values("Date").drop_duplicates(subset="Date")
    df = df.drop(columns=[c for c in ["SNo", "Name", "Symbol"] if c in df.
    ↪columns], errors="ignore")
    cols = [c for c in KEEP_COLS if c in df.columns]
    return df[cols]

df_btc = clean_crypto_df(df_btc_raw)
df_eth = clean_crypto_df(df_eth_raw)

df_btc["btc_ret"] = np.log(df_btc["Close"]).diff()
df_eth["eth_ret"] = np.log(df_eth["Close"]).diff()

df_ret = (
    df_btc[["Date", "btc_ret", "Volume"]]
    .merge(df_eth[["Date", "eth_ret"]], on="Date", how="inner")
    .dropna(subset=["btc_ret", "eth_ret"])
)
df_ret["vol_btc"] = np.log(df_ret["Volume"])

display(df_ret.head())

```

	Date	btc_ret	Volume	eth_ret	vol_btc
1	2015-08-09 23:59:59	0.015534	23,789,600.000000	-0.070710	16.984759
2	2015-08-10 23:59:59	-0.002315	20,979,400.000000	0.009290	16.859052
3	2015-08-11 23:59:59	0.022123	25,433,900.000000	0.410335	17.051593
4	2015-08-12 23:59:59	-0.014942	26,815,400.000000	0.131094	17.104487
5	2015-08-13 23:59:59	-0.008657	27,685,500.000000	0.406292	17.136419

1.2 Section B — Session 1: Distribution & Copula Analysis

All computations below use `btc_ret` and `eth_ret` as required.

```
[7]: def descriptive_stats(series: pd.Series) -> pd.Series:
    return pd.Series({
        "mean": series.mean(),
        "variance": series.var(ddof=1),
        "std": series.std(ddof=1),
        "skewness": series.skew(),
        "kurtosis": series.kurtosis(),
    })

desc = pd.concat(
    {

```

```

        "btc_ret": descriptive_stats(df_ret["btc_ret"]),
        "eth_ret": descriptive_stats(df_ret["eth_ret"]),
    },
    axis=1,
)
cov = df_ret["btc_ret"].cov(df_ret["eth_ret"])
corr = df_ret["btc_ret"].corr(df_ret["eth_ret"])

display(desc)
print(f"Covariance: {cov:.6f}")
print(f"Correlation: {corr:.6f}")

```

```

          btc_ret  eth_ret
mean      0.002259  0.003721
variance   0.001608  0.003856
std       0.040103  0.062099
skewness  -0.821522  0.005249
kurtosis  11.912273  7.667930

```

Covariance: 0.001356
Correlation: 0.544485

```
[8]: def normality_tests(series: pd.Series, label: str):
    mu, sigma = series.mean(), series.std(ddof=1)
    shapiro_stat, shapiro_p = stats.shapiro(series)
    ks_stat, ks_p = stats.kstest(series, "norm", args=(mu, sigma))
    jb_stat, jb_p = stats.jarque_bera(series)
    return {
        "Series": label,
        "Shapiro_W": shapiro_stat,
        "Shapiro_p": shapiro_p,
        "KS_stat": ks_stat,
        "KS_p": ks_p,
        "JB_stat": jb_stat,
        "JB_p": jb_p,
    }

tests = pd.DataFrame([
    normality_tests(df_ret["btc_ret"], "btc_ret"),
    normality_tests(df_ret["eth_ret"], "eth_ret"),
])
display(tests)
```

	Series	Shapiro_W	Shapiro_p	KS_stat	KS_p	JB_stat	JB_p
0	btc_ret	0.898178	0.000000	0.113167	0.000000	12,942.845177	0.000000
1	eth_ret	0.912023	0.000000	0.092244	0.000000	5,261.026821	0.000000

```
[9]: def plot_distribution(series: pd.Series, title: str):
    mu, sigma = series.mean(), series.std(ddof=1)
    x = np.linspace(series.min(), series.max(), 200)
    normal_pdf = stats.norm.pdf(x, mu, sigma)

    fig, axes = plt.subplots(1, 3, figsize=(18, 4))
    sns.histplot(series, bins=50, stat="density", kde=False, ax=axes[0], color="#4C72B0")
    axes[0].plot(x, normal_pdf, color="darkred", lw=2, label="Normal PDF")
    axes[0].set_title(f"Histogram + Normal PDF ({title})")
    axes[0].legend()


    sm.ProbPlot(series, fit=True).qqplot(line="45", ax=axes[1], color="#55A868")
    axes[1].set_title(f"QQ Plot ({title})")

    ecdf = ECDF(series)
    axes[2].plot(ecdf.x, ecdf.y, marker=".", linestyle="none", color="#C44E52")
    axes[2].set_title(f"ECDF ({title})")
    axes[2].set_xlabel("Return")
    axes[2].set_ylabel("F(x)")
    plt.tight_layout()
    plt.show()

plot_distribution(df_ret["btc_ret"], "BTC")
plot_distribution(df_ret["eth_ret"], "ETH")
```


d:\miniconda\envs\ec5303\Lib\site-packages\statsmodels\graphics\gofplots.py:1041: UserWarning: color is redundantly defined by the 'color' keyword argument and the fmt string "b" (-> color=(0.2980392156862745, 0.4470588235294118, 0.6901960784313725, 1)). The keyword argument will take precedence.

```
ax.plot(x, y, fmt, **plot_style)
```



d:\miniconda\envs\ec5303\Lib\site-packages\statsmodels\graphics\gofplots.py:1041: UserWarning: color is redundantly defined by the 'color' keyword argument and the fmt string "b" (-> color=(0.2980392156862745, 0.4470588235294118, 0.6901960784313725, 1)). The keyword argument will take precedence.

```
ax.plot(x, y, fmt, **plot_style)
```




1.2.1 Moment Generating Function (MGF)

Approximate MGF for $t \in [-1, 1]$: $M_X(t) = E[\exp(tX)]$.

```
[10]: def mgf_curve(series: pd.Series, t_values: np.ndarray) -> np.ndarray:
    return np.array([np.exp(t * series).mean() for t in t_values])

t_grid = np.linspace(-1, 1, 200)
mgf_btc = mgf_curve(df_ret["btc_ret"], t_grid)
mgf_eth = mgf_curve(df_ret["eth_ret"], t_grid)

plt.figure(figsize=(8, 5))
plt.plot(t_grid, mgf_btc, label="BTC MGF", color="#4C72B0")
plt.plot(t_grid, mgf_eth, label="ETH MGF", color="#C44E52")
plt.xlabel("t")
plt.ylabel("M_X(t)")
plt.title("MGF Approximation for BTC and ETH Returns")
plt.legend()
plt.grid(True)
plt.show()
```




1.2.2 Probability Integral Transform (PIT)

Empirical CDFs F_{BTC} and F_{ETH} applied to returns to get u_{btc} and v_{eth} in $(0, 1)$.

```
[11]: ecdf_btc = ECDF(df_ret["btc_ret"])
ecdf_eth = ECDF(df_ret["eth_ret"])

df_ret["u_btc"] = ecdf_btc(df_ret["btc_ret"])
df_ret["v_eth"] = ecdf_eth(df_ret["eth_ret"])
df_ret[["u_btc", "v_eth"]] = df_ret[["u_btc", "v_eth"]].clip(1e-6, 1 - 1e-6)

plt.figure(figsize=(6, 5))
plt.scatter(df_ret["u_btc"], df_ret["v_eth"], alpha=0.5, s=10, color="#4C72B0")
plt.xlabel("u_btc = F_BTC(btc_ret)")
plt.ylabel("v_eth = F_ETH(eth_ret)")
plt.title("PIT Scatter: (u_btc, v_eth)")
plt.grid(True)
plt.show()
```



1.2.3 Copula Model Fitting

Fits Gaussian, Student t, Clayton, Gumbel, and Frank copulas on (u_btc, v_eth); computes parameters, log-likelihood, AIC, and BIC.

```
[20]: uv = np.column_stack([df_ret["u_btc"], df_ret["v_eth"]])
n_obs = uv.shape[0]

def resolve_family(candidates, label):
    # Pick the first available BicopFamily attribute from candidates.
    for cand in candidates:
        if hasattr(pv.BicopFamily, cand):
            return getattr(pv.BicopFamily, cand)
    available = [a for a in dir(pv.BicopFamily) if not a.startswith("_")]
    raise AttributeError(
        f"{label} not found in pyvinecopulib.BicopFamily. "
        f" Tried {candidates}. Available: {available}"
    )
```

```

families = {
    "Gaussian Copula": resolve_family(["gaussian"], "Gaussian"),
    "Student t Copula": resolve_family(["t", "student_t", "student"], "Student\u20ac\u20ac\u20ac"),
    "Clayton Copula": resolve_family(["clayton"], "Clayton"),
    "Gumbel Copula": resolve_family(["gumbel"], "Gumbel"),
    "Frank Copula": resolve_family(["frank"], "Frank"),
}

copula_rows = []
for name, family in families.items():
    model = pv.Bicop(family)
    model.fit(uv)
    #     get_parameters   parameters      []
    params = getattr(model, "get_parameters", lambda: getattr(model, \u20ac
    \u20ac"parameters", []))()
    loglik = model.loglik(uv)
    k = len(params)
    aic = 2 * k - 2 * loglik
    bic = np.log(n_obs) * k - 2 * loglik
    copula_rows.append(
        {"Copula": name, "Parameters": params, "LogLik": loglik, "AIC": aic,\u20ac
    \u20ac"BIC": bic}
    )

copula_results = pd.DataFrame(copula_rows).sort_values("AIC").
    \u20acreset_index(drop=True)
best_by_aic = copula_results.loc[copula_results["AIC"].idxmin(), "Copula"]
best_by_bic = copula_results.loc[copula_results["BIC"].idxmin(), "Copula"]

display(copula_results)
print(f"Best copula by AIC: {best_by_aic}")
print(f"Best copula by BIC: {best_by_bic}")

```

	Copula	Parameters	LogLik	\
	AIC	BIC		
0	Student t Copula	[[0.6050088855164852], [2.5998837070963345]]	537.411298	
1	Clayton Copula	[[1.184432884575223]]	461.923295	
2	Frank Copula	[[4.446769639282415]]	426.707511	
3	Gaussian Copula	[[0.5506744917239277]]	391.948074	
4	Gumbel Copula	[[1.5772114867724165]]	372.887469	

	AIC	BIC
0	-1,070.822597	-1,059.467796
1	-921.846591	-916.169190
2	-851.415023	-845.737623
3	-781.896149	-776.218748
4	-743.774938	-738.097538

Best copula by AIC: Student t Copula
 Best copula by BIC: Student t Copula

1.3 Section C — Session 2: Regression Analysis

Dependent variable: eth_ret. Independent variables: btc_ret and vol_btc = ln(Volume_BTC).
 Model 1: ETH on BTC. Model 2: ETH on BTC + BTC volume.

```
[14]: def regression_diagnostics(model, resid_label: str):
    resid = model.resid
    fitted = model.fittedvalues
    fig, axes = plt.subplots(1, 2, figsize=(12, 4))
    sns.scatterplot(x=fitted, y=resid, ax=axes[0], s=12, color="#4C72B0")
    axes[0].axhline(0, color="red", lw=1)
    axes[0].set_title(f"Residuals vs Fitted ({resid_label})")
    sm.qqplot(resid, line="45", ax=axes[1], color="#55A868")
    axes[1].set_title(f"QQ Plot of Residuals ({resid_label})")
    plt.tight_layout()
    plt.show()

reg_data = df_ret.dropna(subset=["btc_ret", "eth_ret", "vol_btc"]).copy()

X1 = sm.add_constant(reg_data[["btc_ret"]])
model1 = sm.OLS(reg_data["eth_ret"], X1).fit()
display(model1.summary())
regression_diagnostics(model1, "Model 1: ETH ~ BTC")
```

Dep. Variable:	eth_ret	R-squared:	0.296			
Model:	OLS	Adj. R-squared:	0.296			
Method:	Least Squares	F-statistic:	908.9			
Date:	Thu, 27 Nov 2025	Prob (F-statistic):	6.27e-167			
Time:	15:57:35	Log-Likelihood:	3316.5			
No. Observations:	2159	AIC:	-6629.			
Df Residuals:	2157	BIC:	-6618.			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	0.0018	0.001	1.618	0.106	-0.000	0.004
btc_ret	0.8431	0.028	30.149	0.000	0.788	0.898
Omnibus:	619.252	Durbin-Watson:	1.843			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	7703.257			
Skew:	0.990	Prob(JB):	0.00			
Kurtosis:	12.039	Cond. No.	24.9			


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

d:\miniconda\envs\ec5303\Lib\site-

```
packages\statsmodels\graphics\gofplots.py:1041: UserWarning: color is
redundantly defined by the 'color' keyword argument and the fmt string "b" (->
color=(0.2980392156862745, 0.4470588235294118, 0.6901960784313725, 1)). The
keyword argument will take precedence.
```

```
    ax.plot(x, y, fmt, **plot_style)
```



```
[15]: X2 = sm.add_constant(reg_data[["btc_ret", "vol_btc"]])
model2 = sm.OLS(reg_data["eth_ret"], X2).fit()
display(model2.summary())
regression_diagnostics(model2, "Model 2: ETH ~ BTC + log(Volume_BTC)")

# Coefficient interpretation
beta1 = model2.params.get("btc_ret", np.nan)
beta2 = model2.params.get("vol_btc", np.nan)
print(f"beta1 (BTC return): {beta1:.6f} → ETH moves beta1 per unit BTC return ↵(ceteris paribus).")
print(f"beta2 (log BTC volume): {beta2:.6f} → ETH return change per unit log ↵volume shock.")
```

Dep. Variable:	eth_ret	R-squared:	0.297			
Model:	OLS	Adj. R-squared:	0.296			
Method:	Least Squares	F-statistic:	455.1			
Date:	Thu, 27 Nov 2025	Prob (F-statistic):	1.30e-165			
Time:	15:57:48	Log-Likelihood:	3317.1			
No. Observations:	2159	AIC:	-6628.			
Df Residuals:	2156	BIC:	-6611.			
Df Model:	2					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	0.0124	0.010	1.276	0.202	-0.007	0.031
btc_ret	0.8428	0.028	30.138	0.000	0.788	0.898
vol_btc	-0.0005	0.000	-1.096	0.273	-0.001	0.000


Omnibus:	600.178	Durbin-Watson:	1.844
Prob(Omnibus):	0.000	Jarque-Bera (JB):	7527.830
Skew:	0.946	Prob(JB):	0.00
Kurtosis:	11.950	Cond. No.	547.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
d:\miniconda\envs\ec5303\Lib\site-
packages\statsmodels\graphics\gofplots.py:1041: UserWarning: color is
redundantly defined by the 'color' keyword argument and the fmt string "b" (->
color=(0.2980392156862745, 0.4470588235294118, 0.6901960784313725, 1)). The
keyword argument will take precedence.
```

```
ax.plot(x, y, fmt, **plot_style)
```



beta1 (BTC return): 0.842828 \rightarrow ETH moves beta1 per unit BTC return (*ceteris paribus*).

beta2 (log BTC volume): -0.000484 \rightarrow ETH return change per unit log volume shock.

1.3.1 Optional Extra Credit

Uncomment and extend below to try nonlinear terms (e.g., `btc_ret^2`) or ARCH/GARCH volatility models using `arch`.

```
[16]: # Example sketch for nonlinear model (not executed by default)
reg_data["btc_ret_sq"] = reg_data["btc_ret"] ** 2
X3 = sm.add_constant(reg_data[["btc_ret", "btc_ret_sq", "vol_btc"]])
model3 = sm.OLS(reg_data["eth_ret"], X3).fit()
display(model3.summary())
```

Dep. Variable:	eth_ret	R-squared:	0.306			
Model:	OLS	Adj. R-squared:	0.305			
Method:	Least Squares	F-statistic:	316.4			
Date:	Thu, 27 Nov 2025	Prob (F-statistic):	3.29e-170			
Time:	15:58:21	Log-Likelihood:	3330.9			
No. Observations:	2159	AIC:	-6654.			
Df Residuals:	2155	BIC:	-6631.			
Df Model:	3					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	0.0091	0.010	0.945	0.345	-0.010	0.028
btc_ret	0.8143	0.028	28.757	0.000	0.759	0.870
btc_ret_sq	-1.0095	0.192	-5.267	0.000	-1.385	-0.634
vol_btc	-0.0003	0.000	-0.583	0.560	-0.001	0.001
Omnibus:	648.525	Durbin-Watson:	1.847			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	7837.661			
Skew:	1.062	Prob(JB):	0.00			
Kurtosis:	12.089	Cond. No.	3.77e+03			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.77e+03. This might indicate that there are strong multicollinearity or other numerical problems.

[]: